

Question 1.2.1

Which class does `LanguageModelingDataset` inherit from?

It inherits from `torch.utils.data.Dataset`

Question 1.2.2

What does the function `lm_collate_fn` do? Explain the structure of the data that results when it is called.

It pads tokens to sentences that are shorter than the longest sentence in the batch. After padding the batch of samples will be composed of sentences of the same length on the specified device.

Question 1.2.3

Looking at the notebook block [6], (with comment "Print out an example of the data") what does this tell you about the relationship between the input (X) and output (Y) that is sent to the model for training?

Y is the latter part of X which is used as the target for the prediction task.

Question 1.2.4

Given one such X,Y pair, how many different training examples does it produce?

3 since there's 4 words in the sentence.

Question 1.2.5

In the `generate` function in the file `model.py` what is the default method for how the generated word is chosen - i.e. based on the model output probabilities?

The default method is to not sample and take the most likely element.

Question 1.2.6

What are the two kinds of heads that model.py can put on to the transformer model? Show (reproduce) all the lines of code that implement this functionality and indicate which method(s) they come from.

Language model heads and classifier heads.

```
self.lm_head = nn.Linear(config.n_embd, config.vocab_size,
bias=False)
self.classifier_head = nn.Linear(config.n_embd,
config.n_classification_class, bias=True)

if not finetune_classify:
    # LM forward procedure
    logits = self.lm_head(x)
else:
    # Finetune classify procedure
    logits = self.classifier_head(x[:, -1, :])
```

They are from the `__init__` and forward method from model.py

Question 1.2.7

How are the word embeddings initialized prior to training?

They are either copied from a pretrained model or randomly initialized depending on whether the type or the parameters are given.

Question 1.2.8

What is the name of the object that contains the positional embeddings?

pos_emb

Question 1.2.9

How are the positional embeddings initialized prior to training?

They are randomly initialized.

Question 1.2.10

Which module and method implement the skip connections in the transformer block? Give the line(s) of code that implement this code.

```
def forward(self, x):  
    x = x + self.mlpf(self.ln_2(x))
```

It'd defined in the forward method.

Question 2.1

Report the value of the loss.

```
iter_dt 0.00ms; iter 0: train loss 10.81249  
iter_dt 8.68ms; iter 100: train loss 5.96995  
iter_dt 9.21ms; iter 200: train loss 2.49559  
iter_dt 10.11ms; iter 300: train loss 1.49280  
iter_dt 15.74ms; iter 400: train loss 0.83902  
iter_dt 8.57ms; iter 500: train loss 0.78918  
iter_dt 9.00ms; iter 600: train loss 0.83952  
iter_dt 9.00ms; iter 700: train loss 0.70429  
iter_dt 9.01ms; iter 800: train loss 0.64494  
iter_dt 9.00ms; iter 900: train loss 0.59071  
iter_dt 9.44ms; iter 1000: train loss 0.56029  
iter_dt 7.98ms; iter 1100: train loss 0.76987  
iter_dt 8.03ms; iter 1200: train loss 0.58646  
iter_dt 9.02ms; iter 1300: train loss 0.61791  
iter_dt 9.00ms; iter 1400: train loss 0.66156  
iter_dt 7.98ms; iter 1500: train loss 0.68874  
iter_dt 8.00ms; iter 1600: train loss 0.69681  
iter_dt 9.00ms; iter 1700: train loss 0.62078  
iter_dt 8.03ms; iter 1800: train loss 0.58298  
iter_dt 9.27ms; iter 1900: train loss 0.59302  
iter_dt 7.99ms; iter 2000: train loss 0.59085  
iter_dt 8.04ms; iter 2100: train loss 0.60756  
iter_dt 9.11ms; iter 2200: train loss 0.64772  
iter_dt 8.00ms; iter 2300: train loss 0.57708  
iter_dt 8.00ms; iter 2400: train loss 0.62975  
iter_dt 9.02ms; iter 2500: train loss 0.65382  
iter_dt 7.99ms; iter 2600: train loss 0.64882  
iter_dt 28.44ms; iter 2700: train loss 0.76117  
iter_dt 22.10ms; iter 2800: train loss 0.63148  
iter_dt 25.01ms; iter 2900: train loss 0.68538
```

Question 2.2

What is the output for each? Why does the the latter parts of the generation not make sense?

'He and I hold the dog.. cat. cat and dog'

'She rubs a cat and dog. cat. cat. cat'

It doesn't make sense as the model was trained on the small corpus where there's no latter part.

Question 2.3

Show the output along with these probabilities for the two examples, and then one of your own choosing.

He 1.0 and 1.0 I 1.0 can 0.7332900166511536 hold 0.6262842416763306 the 0.5019563436508179 dog 0.6777902245521545 . 0.9997881054878235 cat 0.8717184662818909 . 0.9997501969337463 cat 0.7051288485527039 and 0.7348537445068359 cat 0.6792247891426086	She 1.0 rub 1.0 s 1.0 a 0.4766572415828705 dog 0.6089667081832886 and 0.5519123673439026 cat 0.9998014569282532 . 0.9991673231124878 dog 0.6277222633361816 and 0.5781517028808594 cat 0.997319757938385 . 0.999198853969574 cat 0.8431316018104553	The 1.0 cat 1.0 can 1.0 rub 0.6098027229309082 the 0.5069823861122131 dog 0.6066402792930603 . 0.9990702271461487 . 0.9826630353927612 cat 0.6837152242660522 and 0.5920320749282837 cat 0.9993662238121033 and 0.7476220726966858 dog 0.9891071915626526
---	---	---

Question 2.4

Show the result in a table that gives all six words, along with their probabilities, in each column of the table. The number of columns in the table is the total number of generated words. For the first two words generated, explain if the probabilities in the table make sense, given the input corpus.

```
A 1.0
cat 1.0
can 1.0
```

	0	1	2	3	4	5	6	7	8	9
1st	rub 0.52	dog 0.65	a 0.55	a 0.42	a 0.46	a 0.44	. 0.97	s 1.00	. 0.99	a 0.55
2nd	hold 0.48	cat 0.35	the 0.36	the 0.29	the 0.29	the 0.31	and 0.03	and 0.00	. 0.01	the 0.45
3rd	can 0.00	can 0.00	and 0.08	and 0.29	and 0.25	and 0.25	. 0.00	. 0.00	and 0.00	and 0.00
4th	dog 0.00	a 0.00	. 0.00	. 0.00	. 0.00	. 0.00	a 0.00	a 0.00	rub 0.00	holds 0.00
5th	cat 0.00	hold 0.00	s 0.00	cat 0.00	cat 0.00	holds 0.00	the 0.00	cat 0.00	dog 0.00	hold 0.00
6th	a 0.00	rub 0.00	holds 0.00	holds 0.00	holds 0.00	cat 0.00	rub 0.00	dog 0.00	holds 0.00	rub 0.00

Given the input corpus it makes sense since the only possible words after can are rub or hold and they are equally likely. However, dog or cat following rub does not make sense as a determiner is missing in between.

Question 3.2

Report the examples you used and the generation results, and comment on the quality of the sentences.

```
# Example showing how the reloaded model still works
encoded_prompt = train_dataset.tokenizer("Inflation is caused by the").to(trainer.device)
generated_sequence,prob_seq = trainer.model.generate(encoded_prompt, trainer.device, temperature=0.6, max_new_tokens=10)
i=0
import pandas as pd
df=pd.DataFrame({'1st' : [],'2nd' : [],'3rd' : [],'4th' : [],'5th' : [],'6th' : []})
while i<len(prob_seq[0]):...

df.transpose()
Executed at 2023.10.17 23:46:58 in 84ms

In flation is caused by the death
,
F date were to upon

|< < 6 rows > >| 6 rows x 10 columns pd.DataFrame CSV
┌┴┐ 0 ┌┴┐ 1 ┌┴┐ 2 ┌┴┐ 3 ┌┴┐ 4 ┌┴┐ 5 ┌┴┐ 6 ┌┴┐ 7 ┌┴┐ 8
1st death 0.35 \n 0.97 0.94 , 0.45 \n 0.96 F 0.63 date 1.00 were 1.00 to 0.92
2nd ELL 0.15 have 0.02 0.05 . 0.42 0.02 . 0.25 year 0.00 for 0.00 not 0.01
3rd man 0.11 ism 0.00 0.01 and 0.13 the 0.01 ER 0.04 re 0.00 . 0.00 fast 0.01
4th own 0.06 . 0.00 0.00 to 0.00 date 0.00 TER 0.02 of 0.00 0.00 generally 0.01
5th REE 0.04 OF 0.00 ue 0.00 0.00 his 0.00 ; 0.01 had 0.00 ing 0.00 necessarily 0.01

# Example showing how the reloaded model still works
encoded_prompt = train_dataset.tokenizer("The federal government issued").to(trainer.device)
generated_sequence,prob_seq = trainer.model.generate(encoded_prompt, trainer.device, temperature=0.6, max_new_tokens=10)
i=0
import pandas as pd
df=pd.DataFrame({'1st' : [],'2nd' : [],'3rd' : [],'4th' : [],'5th' : [],'6th' : []})
while i<len(prob_seq[0]):...

df.transpose()
Executed at 2023.10.18 00:07:15 in 79ms

The federal government issued in State Secretary
eded of the order years .

|< < 6 rows > >| 6 rows x 10 columns pd.DataFrame CSV
┌┴┐ 0 ┌┴┐ 1 ┌┴┐ 2 ┌┴┐ 3 ┌┴┐ 4 ┌┴┐ 5 ┌┴┐ 6 ┌┴┐ 7 ┌┴┐ 8
1st in 0.82 State 0.37 Secretary 0.52 \n 0.86 eded 0.97 of 1.00 the 0.97 order 0.56 years 0.
2nd from 0.17 report 0.17 ins 0.33 . 0.04 - 0.03 . 0.00 \n 0.03 were 0.13 and 0.19
3rd INT 0.00 position 0.11 that 0.04 ] 0.03 some 0.00 is 0.00 one 0.00 was 0.05 stars 0.
4th , 0.00 commission 0.05 officer 0.02 facing 0.02 af 0.00 - 0.00 it 0.00 ne 0.03 clock 0.
5th of 0.00 States 0.04 of 0.02 0.01 with 0.00 being 0.00 but 0.00 \n 0.03 talents
6th a 0.00 request 0.03 ance 0.01 re 0.00 ual 0.00 ; 0.00 an 0.00 hundred 0.02 pieces 0

# Example showing how the reloaded model still works
encoded_prompt = train_dataset.tokenizer("Spain and Portugal").to(trainer.device)
generated_sequence,prob_seq = trainer.model.generate(encoded_prompt, trainer.device, temperature=0.6, max_new_tokens=10)
i=0
import pandas as pd
df=pd.DataFrame({'1st' : [],'2nd' : [],'3rd' : [],'4th' : [],'5th' : [],'6th' : []})
while i<len(prob_seq[0]):...

df.transpose()
Executed at 2023.10.18 00:09:56 in 86ms

Spain and Portugal coins
coins on laws , vere of were

|< < 6 rows > >| 6 rows x 10 columns pd.DataFrame CSV
┌┴┐ 0 ┌┴┐ 1 ┌┴┐ 2 ┌┴┐ 3 ┌┴┐ 4 ┌┴┐ 5 ┌┴┐ 6 ┌┴┐ 7 ┌┴┐ 8
1st coins 0.45 \n 1.00 coins 0.43 on 0.93 laws 0.56 , 0.81 0.89 vere 0.09 of 1.00
2nd to 0.11 in 0.00 metal 0.21 put 0.04 s 0.15 a 0.05 D 0.05 ule 0.07 the 0.00
3rd has 0.11 and 0.00 coin 0.05 ready 0.01 under 0.13 ., 0.03 C 0.02 OU 0.06 , 0.00
4th for 0.08 at 0.00 gold 0.04 in 0.01 bottom 0.02 and 0.02 \n 0.02 Clerk 0.06 was 0.00
5th a 0.06 , 0.00 metals 0.03 of 0.00 been 0.02 being 0.02 and 0.00 places 0.04 as 0.00
6th out 0.05 the 0.00 are 0.02 more 0.00 _ 0.01 . 0.02 L 0.00 y 0.04 there 0.00
```

The results were gibberish, usually only the first word being predicted will make any sense.

Question 3.3

Report the training and validation curves for the fine-tuning, and the accuracy achieved on the validation dataset.



Question 4.2

Report the classification accuracy on the validation set.
Comment on the performance of this model: is it better than the model you fine-tuned in the previous section?

 [9600/9600 52.44, Epoch 10/10]

Epoch	Training Loss	Validation Loss	Accuracy
1	2.533500	1.490864	0.758333
2	1.413200	1.494483	0.825000
3	0.885100	1.170032	0.812500
4	0.528300	1.163084	0.854167
5	0.380000	1.290416	0.841667
6	0.147600	1.244972	0.833333
7	0.139200	1.334769	0.850000
8	0.090200	1.393531	0.837500
9	0.038800	1.370420	0.870833
10	0.019400	1.279875	0.879167

TrainOutput(global_step=9600, training_loss=0.5977858739693208, metrics={'train_runtime': 3165.1271, 'train_samples_per_second': 3.033, 'train_steps_per_second': 3.033, 'total_flos': 5016897611366400.0, 'train_loss': 0.5977858739693208, 'epoch': 10.0})

The training process was much slower and the validation accuracy got much better, but this is expected since the model used (gpt-2) has 100 times more parameters than gpt-2 nano.